

# Cloud Native Deployments of Bare-Metal High-Performance Al Workflows

#### **ATLAS**

George | Hao | Jing | Shawn | Shubham

MENTOR Chris | Mike | Ravi

#### Introduction:

- Current bare-metal machine in MIT (Satori):
  - 64 1TB memory Power 9 nodes.
  - Each node hosts four NVidia V100 32GB memory GPU cards.
  - 100Gb/s Infiniband network connects the cluster nodes.
- OpenShift/Kubernetes migration
- AI workflow between bare-metal and OpenShift
- Potential benefits in OpenShift

#### Vision and Goals:

- Survey existing MIT bare-metal workloads & containerize one of them.
- Monitor & compare OpenShift workflows and bare-metal workflows.
- Generate reports with pros/cons of migrating bare-metal workflow to OpenShift environment.

### Users/Personas:

#### This system will target the following users:

• Al researchers looking to deploy high-performance workflows:

bare-metal environment  $\rightarrow$  cloud native environment.

- ML/Al engineers looking to deploy processing-hungry system to cloud.
- Users seeking data privacy in the cloud (provided by OpenShift).
- Users needing Singularity/virtualization systems to containerize workflows in HPC (high-performance computing) clusters.

#### Users/Personas:

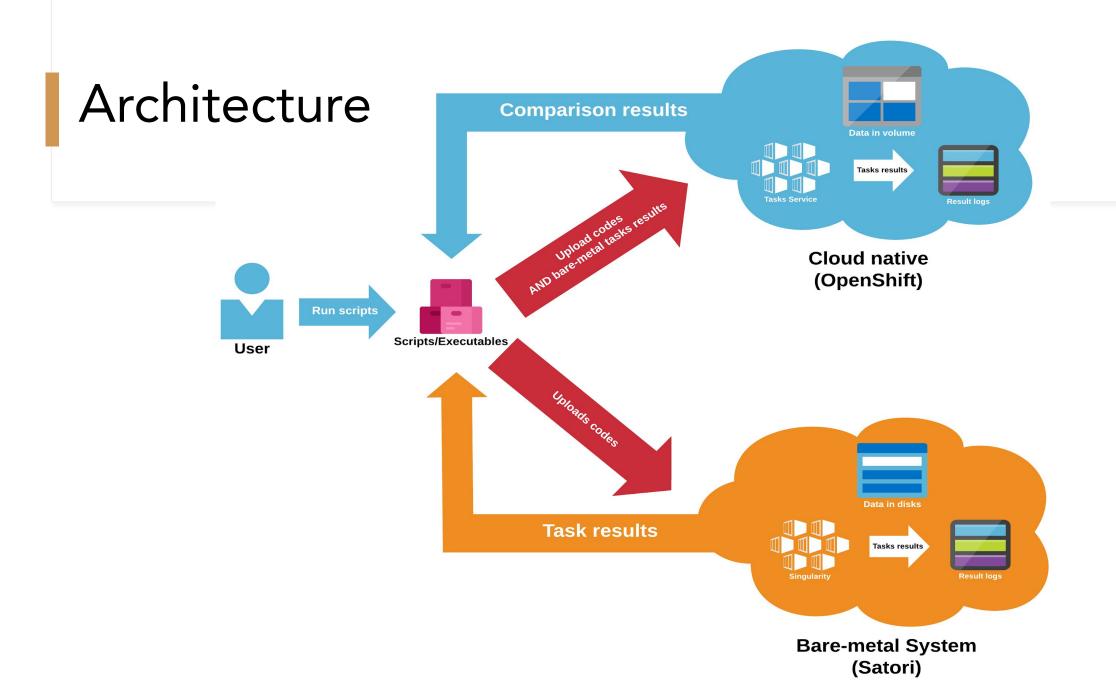
- A quintessential example of a user could be the MIT-IBM Watson AI laboratory looking to scale their workflows into the cloud in a discrete fashion.
- Average users/hobbyists looking to deploy non-intensive computational processes to the cloud.

#### This system will NOT target the following users:

• Users with complex requirements who might require additional interface/systemic modification.

## Scope and Features:

- Create documentation & scripts to allow users to containerize existing
   High Performance (AI) workflows
- Generate charting comparing performance metrics (elasticity, economics, etc.) between bare-metal & OpenShift environments.
- Generate display (of suggestions) for 'under-utilized' nodes in OpenShift (*running backfill workloads*?).
- Deploy researcher workflows/code with ease from bare-metal environment to OpenShift/Kubernetes.



## Acceptance Criteria:

- Interface: containerize & deploy a specific AI workflow.
- Generate comparisons (elasticity, performance, economics, etc.)
   between project being run in OpenShift vs Satori.
- Directing resources to under-utilized nodes (or displaying there are such instances).
- Extending to wider class of projects by circumventing problem of workflows being tied to a current system.

# BigGAN - MIT HPC

- Generative Adversarial Network
- Novel/creative data
- Realistic: Inception Score (IS) of 166.5



#### Thus far...Satori

- BigGAN (138Gb training data)
- Submitting jobs to queue
- WMLCE Environment -- PowerAl

```
(wmlce-1.6.2) [sorescu@service0001 BigGAN-PyTorch]$ bsub < jobs/biggan_deep128_imagenet.lsf
Job <24334> is submitted to queue <normal>.
(wmlce-1.6.2) [sorescu@service0001 BigGAN-PyTorch]$ ls
animal hash.py
                               datasets.py
                                                                  notebooks
                                                                               SATORI.md
                                                   layers
BigGANdeep.py
                               eval models.py
                                                    layers.py
                                                                  profiler.py scripts
BigGAN.py
                               get power
                                                    LICENSE
                                                                   pycache
                                                                               set.sh
```

# Satori Logs

- Iterative training plots
- Tracking in JupyterNotebooks

#### Plot Generator and Discriminator loss trajectories for a particular experiment

```
log_dirs = [os.path.join('logs', f) for f in os.listdir('logs') if os.path.isdir(os.path.join('logs', f))]
for i, lg in enumerate(log_dirs):
    print(i, lg)

log_num = 0
logs = utils.load_logs(log_dirs[log_num])

utils.plot_loss_logs(logs['loss'], smoothing=0.000, figsize=(24, 12))
```

- 0 logs/BigGAN\_ImageNet\_128\_biggan\_deep\_seed0\_Gch128\_Dch128\_Gd2\_Dd2\_bs256\_nDs2\_Glr5.0e-05\_Dlr2.0e-04\_Gnlinplace\_relu\_Dnlinplace\_relu\_Ginitortho\_Dinitortho\_Gattn64\_Dattn64\_cr\_Gshared\_hier\_ema
- 1 logs/BigGAN\_ImageNet\_128\_biggan\_deep\_seed0\_Gch128\_Dch128\_Gd2\_Dd2\_bs256\_nDs2\_Glr5.0e-05\_Dlr2.0e-04\_Gnlinplace\_relu\_Dnlinplace\_relu\_Ginitortho\_Dinitortho Gattn64 Dattn64 Gshared hier ema



# Satori Logs

- Metadata logs
- Track performance: bare-metal
   vs cloud

```
datetime: 2020-02-10 10:16:34.361624
config: {'dataset': 'ImageNet', 'resolution': 128, 'augment': False, 'num_workers': 2, 'pi
'use multiepoch sampler': False, 'world size': 32, 'rank': 0, 'dist url': 'file:///nobacku
'dist backend': 'nccl', 'multiprocessing distributed': True, 'gpu': 0, 'dataset type': 'Im
'D param': 'SN', 'G ch': 128, 'D ch': 128, 'G depth': 2, 'D depth': 2, 'D wide': True, 'G
'hier': True, 'cross replica': False, 'mybn': False, 'G nl': 'inplace relu', 'D nl': 'inpl
'bn', 'seed': 0, 'G init': 'ortho', 'D init': 'ortho', 'skip init': False, 'G lr': 5e-05,
'D B2': 0.999, 'batch size': 64, 'G batch size': 0, 'num G accumulations': 1, 'num D steps
'num epochs': 500, 'parallel': False, 'G_fp16': False, 'D_fp16': False, 'D_mixed_precision
False, 'num standing accumulations': 16, 'use torch FID': False, 'pretrained': 'imagenet',
2, 'num best copies': 5, 'which best': 'IS', 'no fid': False, 'test every': 2000, 'num inc
'data root': 'data', 'weights root': 'weights', 'logs root': 'logs', 'samples root': 'samp
'name suffix': '', 'experiment name': '', 'config from name': False, 'ema': True, 'ema dec
1e-06, 'BN eps': 1e-05, 'SN eps': 1e-06, 'num G SVs': 1, 'num D SVs': 1, 'num G SV itrs':
'toggle grads': True, 'which train fn': 'GAN', 'load weights': '', 'resume': False, 'logst
False, 'sv_log_interval': 10, 'debug': False, 'n_classes': 1000, 'G_activation': ReLU(inpl
True}
```

# Containerize Big-GAN workflow

- 1. Write bash scripts to automate the entire workflow (data preprocess -> pre-fetch eval metric scores -> model training)
- 2. Build a docker image that runs the workflow in a container
- 3. Run & monitor the container (redirect std out to log file)

```
AN$ docker build -t big-gan .
ending build context to Docker daemon 7.68kB
Step 1/9 : FROM centos:centos7
Step 2/9: RUN yum -y update; yum clean all
 ---> Using cache
 ---> 52a0c208bf19
Step 3/9 : RUN yum -y install epel-release gcc openssl-devel bzip2-de
vel git python37 python3-pip vim; yum clean all
 ---> Usina cache
 ---> b09d7e2135cd
Step 4/9 : ADD ./requirements.txt /python/requirements.txt
 ---> Usina cache
 ---> e6f316b93acc
 tep 5/9 : RUN pip3 install --upgrade pip && pip install -r /python/r
equirements.txt
 ---> Using cache
 ---> 220ce5b216b7
Step 6/9 : ADD ./ /BigGAN
 ---> Using cache
 ---> 686eafbfc818
tep 7/9 : WORKDIR /BigGAN
 ---> Using cache
 ---> ced3018731d2
 tep 8/9 : CMD cd /BigGAN
 ---> Using cache
 ---> 182f9585b109
Step 9/9 : CMD /bin/sh run.sh
 ---> Using cache
 ---> 798971a0cd44
uccessfully built 798971a0cd44
Successfully tagged big-gan:latest
```

# Containerize Big-GAN workflow - stdout logs

```
ws/workflows/BiaGAN$ docker loas -f 5cc518cf0b84
Cloning into 'BigGAN-PyTorch'...
Switched to a new branch 'satori'
Branch satori set un to track remote branch satori from origin
                                                                                                          Cloning code from github repo
Making hdt5...
 100%|#########| 1300/1300 [00:00<00:00, 112733.79it/s]
 100%|#########| 1/1 [00:00<00:00, 11.11it/s]
 {'dataset': 'ImageNet', 'resolution': 128, 'data_root': '/data/', 'batch_size': 256,
Generating index file ImageNet_train.npz
                                                                                                                 Data preprocess & compress
Saving new hdf5 file to : /data/ImageNet/ImageNet-128.hdf5
HDF5 file /data/ImageNet/ImageNet-128.hdf5 already exists!
 alculating inception moments...
 Downloading: "https://download.pytorch.org/models/inception_v3_google-1a9a5a14.pth'
 L00%|######### 104M/104M [00:06<00:00, 16.3MB/s]
 ['dataset': 'ImageNet', 'resolution': 128, 'dataset_type': 'ImageFolder', <u>'data</u>root
 False, 'num_workers': 1, 'shuffle': False, 'seed': 0, 'pretrained': 'imagenet'}
 _oading pre-saved index file ImageNet_train.npz
 oading inception net...
 oading inception model in cpu mem...
                                                                                                                    Pre-fetch eval metric
 Getting logits from pretrained model...
 L00%|#########| 325/325 [03:59<00:00, 1.36it/s]
 Calculating inception metrics...
 imagenet model evaluated on ImageNet-128 has IS of 2.16672 +/- 0.17320
 Calculating means and covariances...
 saving calculated means and covariances to disk...
 Start training bigGAN...
 Making directory weights for weights_root...
 Making directory logs for logs_root...
 Making directory samples for samples_root...
                                                                                                                    Model training
 Making directory results for results_root...
Param count for Ds initialized parameters: 34590209
 Initializing EMA parameters to be source parameters...
 oading pre-saved index file ImageNet_train.npz
```

oading inception model in cpu mem...

#### Demo

- 1. git clone https://github.com/BU-CLOUD-S20/Cloud-native-deployments-of-bare-metal-high-performance-Al-workflows.git
- 2. cd Cloud-native-deployments-of-bare-metal-high-performance-Al-workflows/workflows/BigGAN/
- 3. docker build -t big-gan .
- 4. docker run -d --name cs528\_demo -v /Users/Shawn/Documents/BU/courses/2020\_spring/CComp/Cloud-native-deployments-of-bare-metal-high-performance-Al-workflows/data/:/d ata/ big-gan:latest
- 5. docker exec -it 5cc518cf0b84 /bin/bash

## Optional Features:

- Automatic deploy: interface or containers to automatically execute experimental codes in two different environments.
- Generalized: orientation towards high-performance AI workflows, but have the capability to deploy wide range of projects.
- Operate with ease across multiple deployments (MIT HPC, MIT-IBM Watson lab, etc.).

# Thank you!

Q&A